

Bunyan Logging in Production at Joyent

Trent Mick (@trentmick)

Joyent

Vancouver Node.js Meetup, 7 Aug 2014

Bunyan Logging ...

- Blow through usage and some features.

... in production at Joyent

- Show a few (hopefully) interesting examples from production experiences.
- Discuss some motivations.

Bunyan is

① A node.js lib for structured service (and CLI) logging

```
$ npm install bunyan  
$ cat foo.js
```

```
var bunyan = require('bunyan');  
var log = bunyan.createLogger({  
  name: 'foo',  
  // ...  
});
```

```
log.info('hi');
```

```
var err = new Error('ack');  
log.error(err, 'boom!');
```

Bunyan is

② A (lightweight) JSON log format

- one line of JSON per record, small set of minimal fields:

```
{ "name": "foo", "hostname": "grape.local", "pid": 68313, "level": 30,  
  "msg": "hi", "time": "2014-08-07T06:18:53.057Z", "v": 0 }
```

- The primary audience of logs is machine processing.

Bunyan is

③ A bunyan CLI for viewing and filtering logs

Reading JSON sucks, but ANSI codes in files + custom parsing code suck more:

```
$ node foo.js | bunyan
[2014-08-07T06:35:48.340Z]  INFO: foo/68813 on grape.local: hi
[2014-08-07T06:35:48.343Z] ERROR: foo/68813 on grape.local: boom!
    Error: ack
      at Object.<anonymous> (/Users/trentm/tmp/play/foo.js:9:11)
      at Module._compile (module.js:456:26)
      at Object.Module._extensions..js (module.js:474:10)
```

```
bunyan foo.log [bar.log ...]
bunyan -l warn foo.log
bunyan -c 'this.user=="bob"' foo.log
```

Bunyan features

- standard *levels*: trace, debug, info, warn, error, fatal
- extensible *streams* for directing logs to stderr, stdout, files, rotating files, custom endpoints
- specialization of Loggers with `log.child`
- custom rendering of logged objects with *serializers*
- runtime log snooping via DTrace support

See [README](#) for more.

Joyent[®] **SmartDataCenter 7**

- ~40 service/agent/tool log types (~90% bunyan)
- 1-5 service instances, one agent instance per server
- ~12 DCs (counting test and staging DCs)

Motivation: Log more details

- Make it convenient in code to log more and appropriate details.
- Painful, limited (traditional):

```
log.info('%s %s %s %s', req.method, req.url,  
        res.statusCode, res.latency)
```

- Easier to code (bunyan):

```
log.info({req: req}, 'received request');  
log.debug({res: res}, 'responded');  
log.error({err: err, path: path}, 'save failed');
```


Log more details: render nicely

Make having the extra details not a burden when browsing logs by rendering them nicely.

HTTP request

```
[2014-08-07T17:08:24.071Z] INFO: imgapi/74427 on e540e25b-6afc-41ee-a940-a1da39628676: handled: 200
GET /images/2b683a82-a066-11e3-97ab-2faa44701c5a HTTP/1.1
accept: application/json
user-agent: restify/2.8.1 (ia32-sunos; v8/3.14.5.9; OpenSSL/1.0.1h) node/0.10.29
accept-version: *
date: Thu, 07 Aug 2014 17:08:22 GMT
host: imgapi.nightly-1.joyent.us
connection: close
```

HTTP response

```
[2014-08-07T17:12:01.601Z] INFO: cnapi/12183 on 7b44b026-e23d-4117-9093-9c37da632501: handled: 500
...
--
HTTP/1.1 500 Internal Server Error
content-type: application/json
content-length: 77
date: Thu, 07 Aug 2014 17:12:01 GMT
server: Compute Node API
x-request-id: f2692cf0-1e55-11e4-a698-b90f8ef63a7a
x-response-time: 1
x-server-name: 7b44b026-e23d-4117-9093-9c37da632501
```

HTTP response + internal error details

```
[2014-08-07T17:12:01.601Z] INFO: cnapi/12183 on 7b44b026-e23d-4117-9093-9c37da632501: handled: 500
...
--
HTTP/1.1 500 Internal Server Error
content-type: application/json
content-length: 77
date: Thu, 07 Aug 2014 17:12:01 GMT
server: Compute Node API
x-request-id: f2692cf0-1e55-11e4-a698-b90f8ef63a7a
x-response-time: 1
x-server-name: 7b44b026-e23d-4117-9093-9c37da632501
--
InternalServerError: Precondition failed: no connection moray
  at Object.ensureConnectedToMoray [as moray] (/opt/smartdc/cnapi/lib/endpoints/index.js:38:14)
  at Server.<anonymous> (/opt/smartdc/cnapi/lib/endpoints/index.js:163:39)
  at next (/opt/smartdc/cnapi/node_modules/restify/lib/server.js:736:42)
  at f (/opt/smartdc/cnapi/node_modules/once/once.js:16:25)
  at ensureConnectionTimeout (/opt/smartdc/cnapi/lib/endpoints/index.js:65:5)
  at Server.<anonymous> (/opt/smartdc/cnapi/lib/endpoints/index.js:149:9)
  at next (/opt/smartdc/cnapi/node_modules/restify/lib/server.js:736:42)
  at f (/opt/smartdc/cnapi/node_modules/once/once.js:16:25)
  at Server.<anonymous> (/opt/smartdc/cnapi/lib/endpoints/index.js:145:9)
```

Error + underlying error cause

```
[2014-08-07T17:15:28.452Z] INFO: imgapi/74427 on e540e25b-6afc-41ee-a940-a1da39628676: handled: 404
...
ResourceNotFoundError: image not found; caused by ObjectNotFoundError: imgapi_images::f3785e6b-cb9b-
  at /opt/smartdc/imgapi/lib/database.js:546:26
  at f (/opt/smartdc/imgapi/node_modules/once/once.js:16:25)
  at EventEmitter.<anonymous> (/opt/smartdc/imgapi/node_modules/moray/lib/objects.js:137:9)
  at EventEmitter.g (events.js:180:16)
  at EventEmitter.emit (events.js:95:17)
  at Client._handleMessage (/opt/smartdc/imgapi/node_modules/moray/node_modules/fast/lib/client.js
  at MessageDecoder.onMessage (/opt/smartdc/imgapi/node_modules/moray/node_modules/fast/lib/client
  at MessageDecoder.EventEmitter.emit (events.js:95:17)
  at MessageDecoder._write (/opt/smartdc/imgapi/node_modules/moray/node_modules/fast/lib/protocol/
  at doWrite (/opt/smartdc/imgapi/node_modules/moray/node_modules/fast/node_modules/readable-strea
Caused by: ObjectNotFoundError: imgapi_images::f3785e6b-cb9b-e31b-8ac5-852e57d2662a does not exist
  at EventEmitter.<anonymous> (/opt/smartdc/moray/lib/objects/get.js:102:16)
  at EventEmitter.g (events.js:180:16)
  at EventEmitter.emit (events.js:95:17)
  at endOnError (/opt/smartdc/moray/lib/pg.js:89:13)
  at f (/opt/smartdc/moray/node_modules/once/once.js:16:25)
  at onQueryEnd (/opt/smartdc/moray/lib/pg.js:116:9)
  at EventEmitter.emit (events.js:95:17)
  at Query.handleReadyForQuery (/opt/smartdc/moray/node_modules/pg/lib/query.js:86:8)
  at null.<anonymous> (/opt/smartdc/moray/node_modules/pg/lib/client.js:159:19)
  at EventEmitter.emit (events.js:117:20)
  at null.<anonymous> (/opt/smartdc/moray/node_modules/pg/lib/connection.js:97:12)
```

Maintain error cause chains with [node-verror](#). Bunyan will render these.

Log more details: filtering

Motivate logging of more structured details by making filtering convenient and compelling.

```
# ERROR level and above  
$ bunyan foo.log -l error
```

```
# Watch incoming requests.  
$ tail -f foo.log | bunyan -c 'this.req'
```

```
# Server errors in HTTP responses.  
$ tail -f foo.log | bunyan -c 'this.res && this.res.statusCode >= 500'
```


Log more details: serializers

- Problem: Dumping full objects -> extraneous goop in the logs.
- Define how an obj is logged with serializer functions for well-defined log *field names*.
- Bunyan defines for `err`, `req`, and `res`. Use is explicit, e.g.:

```
var log = bunyan.createLogger({
  //...
  serializers: {
    err: bunyan.stdSerializers.err,
    // custom
    user: function (u) { if (u) return {login: u.login, name: u.name}; }
  }
});

log.info({user: theUser}, 'logged in');
log.error({err: e, user: theUser}, 'logged failed');
```



```
[2014-08-07T17:25:55.126Z] INFO: imgapi/74427 on e540e25b-6afc-41ee-a940-a1da396
POST /images?action=create-from-vm&vm_uuid=c691c60d-8f65-6eb2-81d6-d163d65ea4
x-request-id: dd8d4ad0-1e57-11e4-8499-df7725ba1cdc
accept: application/json
content-type: application/json
user-agent: restify/2.6.1 (ia32-sunos; v8/3.14.5.9; OpenSSL/1.0.1h) node/0.10
date: Thu, 07 Aug 2014 17:25:48 GMT
content-length: 65
content-md5: d02bUF89wkeSJT+jGTEENg==
host: imgapi.nightly-1.joyent.us
connection: keep-alive
```

```
{
  "name": "1cfefb76f-d951-c1e3-b9e7-97825c7c4efc",
  "version": "1.0.0",
  "owner": "551a64ca-3e08-c0de-cfd2-da81641890ec",
  "uuid": "1c7a2b2d-f5da-eed2-d37c-e3a01de9a8ee",
  "requirements": {
    "min_platform": {
      "7.0": "20140807T022845Z"
    }
  }
}
```

```
--
```

```
HTTP/1.1 200 OK
```

Necessity: balancing volume

- Serializers help limit the size of a logged object.
- Obviously, set runtime log level to control log volume.
- Not enough.
- Two interesting examples: RequestCaptureStream and DTrace.

Balancing volume: RequestCaptureStream

- Provided by restify web service library.
- Store low level log records in in-memory ring buffer, *per request id*.
- If there is a `log.warn` or above, then all low level records for *that request id* are logged.
- Used in Joyent's public web APIs.

Code for a RequestCaptureStream

```
var log = bunyan.createLogger({
  name: 'cloudapi',
  serializers: restify.bunyan.serializers,
  streams: [
    {
      level: 'info',
      stream: process.stderr
    },
    {
      type: 'raw',
      stream: new RequestCaptureStream({
        level: bunyan.WARN,
        maxRecords: 100,
        maxRequestIds: 100,
        streams: [ {
          stream: process.stderr
        } ]
      })
    }
  ]
});
```

Balancing volume: DTrace

- Typically log level is INFO or DEBUG at best.
- How to get TRACE logging details?
- Option 1: Update app config and restart.
- Option 2: DTrace bunyan provider. No restart.

```
$ bunyan -p $PID
$ bunyan -p $NAME      # search processes by name via `pgrep` or `ps`
$ bunyan -p '*'        # instrument *all* processes!
$ bunyan -p '*' -l DEBUG
$ bunyan -p '*' -c 'this.name === "cloudapi"'
```

`bunyan -p` demo

Motivation: Tie things together

- Service logs can be a jumble of simultaneous operations.
- How to meaningfully separate?

Tie things together: module-scope loggers don't help

```
lib/main.js:    var log = log4j.getLogger('foo')  
lib/db.js:    var log = log4j.getLogger('foo.db')  
lib/model.js:    var log = log4j.getLogger('foo.model')
```

Common with log4j-inspired libraries.
Groups on module, not on related operations.

Tie things together: always log context -> PITA

```
log.info('user=%s start blerping', user)
log.info('user=%s frob the frib', user);
log.info('bling bloop');    // oops, forgot to log `user`
```

- Easy to forget to include relevant context.
- Labourious to add, so gets skipped.

Tie things together: `log.child`

```
req.log = app.log.child({req_id: req.getId()});  
// ...  
req.log.info({route: route}, 'routed request');  
// ...  
req.log.debug(err, 'validation failed');
```

- Group related logging via logger specialization with `log.child`.
- Annoyance: Now need to pass a logger around code.
- Attach to relevant context object, if available.
E.g. `restify` creates a `<request>.log` for each HTTP request.

Tie things together: req_id

```
$ bunyan foo.log \  
  -c 'this.req_id === 'd6740300-1e58-11e4-8f85-2fef7c7c36dd''
```

Now we can see all logging for a particular request.

Tie things together: `req_id++`

How about tracking a request through multiple services?
Example: provision a VM in SmartDataCenter.

- Cloud API: public facing API
- VM API: coordinates creation of VM
- Server API (CNAPI): handles talking to server agents
- Network API (NAPI): network provisioning
- Firewall API (FWAPI): firewall management for VMs

Tie things together: `req_id++`

Demo tracking VM creation in SmartDataCenter via `req_id`.

Questions?

Trent Mick (@trentmick)

github.com/trentm/node-bunyan
bunyan-logging@googlegroups.com

Slides: trentm.com/talk-bunyan-in-prod/
and on GitHub: github.com/trentm/talk-bunyan-in-prod